

**Problem 1: The salmonella data set from the faraway library was collected in a salmonella reverse mutagenicity assay where the numbers of revertant colonies of TA98 Salmonella observed on each of three replicate plates for different doses of quinoline**

**(a) Fit a linear model with colonies as the response and log(dose + 1) as a predictor.**

**(b) Make residual plots and comment on the results.**

**(c) Check this model for lack of fit using an appropriate test.**

```
In [22]: library('faraway')
library('ggplot2')
options(repr.plot.width=6, repr.plot.height=4)
data(salmonella)
```

```
In [193]: head(salmonella, 2)
```

| colonies | dose | residuals |
|----------|------|-----------|
| 15       | 0    | -4.823482 |
| 21       | 0    | 1.176518  |

**(a) Fit a linear model with colonies as the response and log(dose + 1) as a predictor.**

```
In [201]: model = lm(colonies ~ log(dose + 1), data = salmonella)
```

```
In [202]: summary(model)
```

Call:

```
lm(formula = colonies ~ log(dose + 1), data = salmonella)
```

Residuals:

| Min     | 1Q     | Median | 3Q    | Max    |
|---------|--------|--------|-------|--------|
| -16.376 | -6.882 | -1.509 | 5.400 | 29.119 |

Coefficients:

|               | Estimate | Std. Error | t value | Pr(> t )   |
|---------------|----------|------------|---------|------------|
| (Intercept)   | 19.823   | 5.064      | 3.915   | 0.00123 ** |
| log(dose + 1) | 2.396    | 1.128      | 2.125   | 0.04955 *  |

---

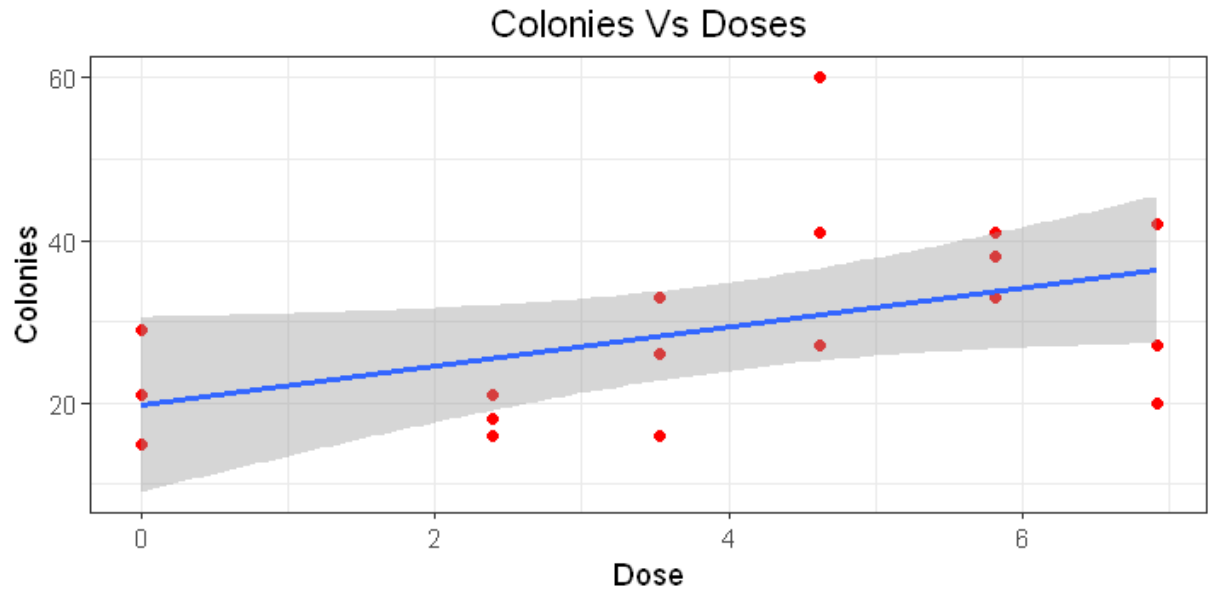
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 10.84 on 16 degrees of freedom

Multiple R-squared: 0.2201, Adjusted R-squared: 0.1713

F-statistic: 4.514 on 1 and 16 DF, p-value: 0.04955

```
In [203]: ggplot(data = salmonella, aes(x = log(salmonella$dose+1), y = salmonella$colonies)) +
  geom_point() +
  theme_update(plot.title = element_text(hjust = 0.5)) +
  theme_set(theme_bw()) +
  geom_point(color='red') +
  labs(title='Colonies Vs Doses', x='Dose', y = 'Colonies') +
  geom_smooth(method = lm)
```



**(b) Make residual plots and comment on the results.**

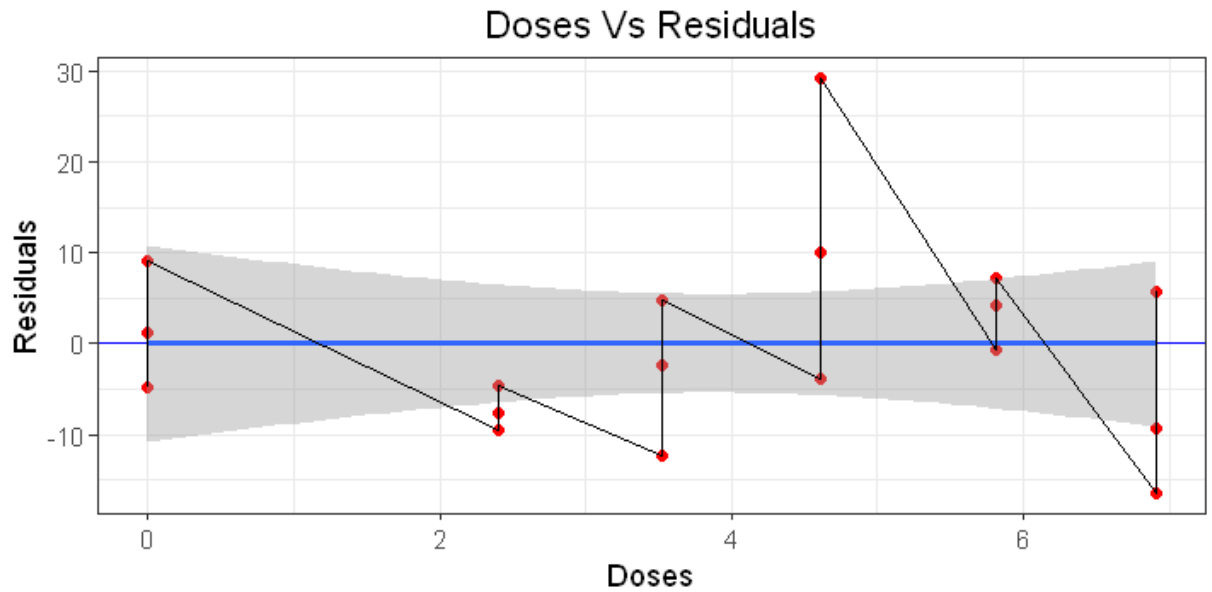
```
In [204]: salmonella$residuals = model$res
```

```
In [205]: head(salmonella)
```

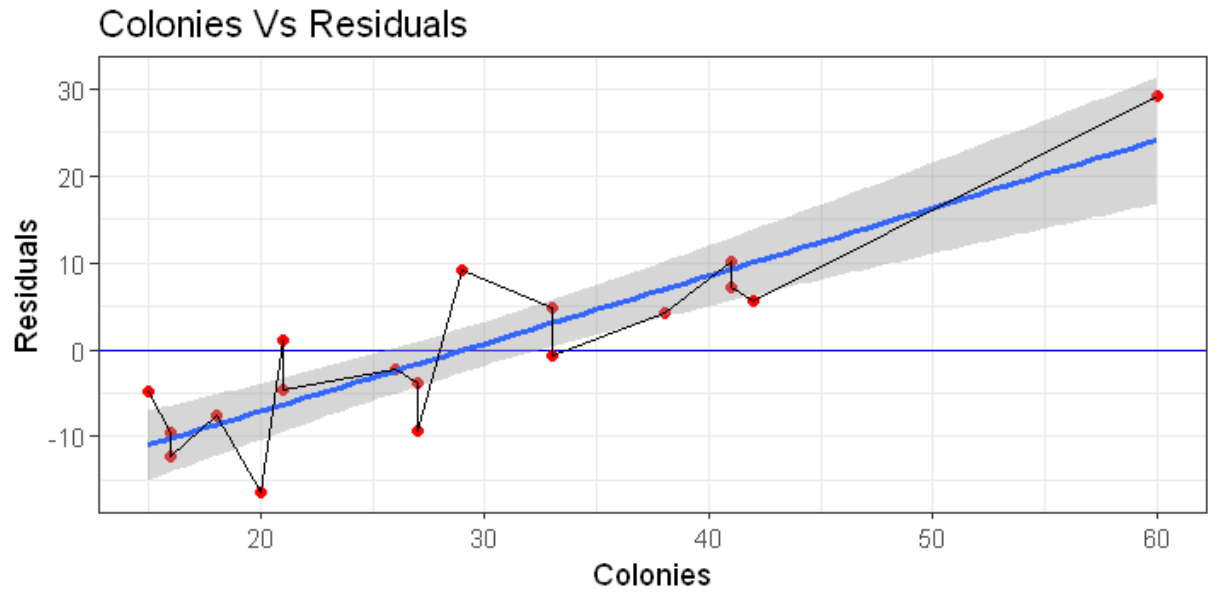
| colonies | dose | residuals |
|----------|------|-----------|
| 15       | 0    | -4.823482 |
| 21       | 0    | 1.176518  |
| 29       | 0    | 9.176518  |
| 16       | 10   | -9.568484 |
| 18       | 10   | -7.568484 |
| 21       | 10   | -4.568484 |

```
In [206]: library('ggplot2')
options(repr.plot.width=6, repr.plot.height=3)


ggplot(data = salmonella, aes(x = log(salmonella$dose+1), y = residuals))+
  geom_point()+
  geom_hline(yintercept = 0, color = 'blue') +
  theme_update(plot.title = element_text(hjust = 0.5))+
  theme_set(theme_bw())+
  geom_point(color='red')+
  labs(title='Doses Vs Residuals', x='Doses', y = 'Residuals')+
  geom_smooth(method = lm)+
  geom_line()
```



```
In [207]: ggplot(data = salmonella, aes(x = salmonella$colonies, y = residuals))+  
  geom_point()+  
  geom_hline(yintercept = 0, color = 'blue') +  
  geom_point(color='red')+  
  labs(title='Colonies Vs Residuals', x='Colonies', y = 'Residuals')+  
  geom_smooth(method = lm)+  
  geom_line()
```



```
In [209]: ggplot(data = salmonella, aes(x = model$fitted, y = residuals))+
  geom_point()+
  geom_hline(yintercept = 0, color = 'blue') +
  geom_point(color='red')+
  labs(title='fitted Vs Residuals model', x='fitted', y = 'Residuals')+
  geom_smooth(method = lm)+
  geom_line()
```



- When we make residuals vs doses plot, we can see that there is no clear relationship that emerges.
- When we make residuals vs colonies plot, linear relationship is seen.
- When we make residuals vs fitted plot, no clear relationship emerges. To confirm, we can perform Durbin-Watson test.

```
In [45]: # install.packages('lmtest')
# library('lmtest')
dwtest(model)
```

#### Durbin-Watson test

```
data: model
DW = 1.6279, p-value = 0.1382
alternative hypothesis: true autocorrelation is greater than 0
```

- The p value is greater than 5% significance level, hence we accept null hypothesis. i.e. residuals are not correlated.

#### (c) Check this model for lack of fit using an appropriate test.

- Since we do not know the sigma value of each  $X_i$ , hence we need to check lack of fit when sigma unknown


```
In [46]: head(salmonella, 2)
```

| colonies | dose | residuals |
|----------|------|-----------|
| 15       | 0    | -4.823482 |
| 21       | 0    | 1.176518  |

```
In [47]: salmonella[order(salmonella$colonies),]
```

|    | colonies | dose | residuals   |
|----|----------|------|-------------|
| 1  | 15       | 0    | -4.8234823  |
| 4  | 16       | 10   | -9.5684844  |
| 7  | 16       | 33   | -12.2721201 |
| 5  | 18       | 10   | -7.5684844  |
| 16 | 20       | 1000 | -16.3758360 |
| 2  | 21       | 0    | 1.1765177   |
| 6  | 21       | 10   | -4.5684844  |
| 8  | 26       | 33   | -2.2721201  |
| 10 | 27       | 100  | -3.8806279  |
| 17 | 27       | 1000 | -9.3758360  |
| 3  | 29       | 0    | 9.1765177   |
| 9  | 33       | 33   | 4.7278799   |
| 13 | 33       | 333  | -0.7461159  |
| 14 | 38       | 333  | 4.2538841   |
| 11 | 41       | 100  | 10.1193721  |
| 15 | 41       | 333  | 7.2538841   |
| 18 | 42       | 1000 | 5.6241640   |
| 12 | 60       | 100  | 29.1193721  |

```
In [89]: model_a = lm(colonies ~ log(dose + 1), data = salmonella);  
plot(colonies ~ log(dose + 1), data=salmonella); abline(coef(model_a));
```



```
In [90]: model_b = lm(colonies ~ factor(log(dose + 1)), data = salmonella);
cbind(salmonella, model_b$fitted)[order(salmonella$colonies),]
```

|    | colonies | dose | residuals   | model_b\$fitted |
|----|----------|------|-------------|-----------------|
| 1  | 15       | 0    | -4.8234823  | 21.66667        |
| 4  | 16       | 10   | -9.5684844  | 18.33333        |
| 7  | 16       | 33   | -12.2721201 | 25.00000        |
| 5  | 18       | 10   | -7.5684844  | 18.33333        |
| 16 | 20       | 1000 | -16.3758360 | 29.66667        |
| 2  | 21       | 0    | 1.1765177   | 21.66667        |
| 6  | 21       | 10   | -4.5684844  | 18.33333        |
| 8  | 26       | 33   | -2.2721201  | 25.00000        |
| 10 | 27       | 100  | -3.8806279  | 42.66667        |
| 17 | 27       | 1000 | -9.3758360  | 29.66667        |
| 3  | 29       | 0    | 9.1765177   | 21.66667        |
| 9  | 33       | 33   | 4.7278799   | 25.00000        |
| 13 | 33       | 333  | -0.7461159  | 37.33333        |
| 14 | 38       | 333  | 4.2538841   | 37.33333        |
| 11 | 41       | 100  | 10.1193721  | 42.66667        |
| 15 | 41       | 333  | 7.2538841   | 37.33333        |
| 18 | 42       | 1000 | 5.6241640   | 29.66667        |
| 12 | 60       | 100  | 29.1193721  | 42.66667        |



```
In [91]: summary(model_b)
```

```
Call:
lm(formula = colonies ~ factor(log(dose + 1)), data = salmonella)

Residuals:
    Min       1Q   Median       3Q      Max
-15.667  -3.917  -0.500   3.417  17.333

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)      21.667      5.506   3.935  0.00198 **
factor(log(dose + 1))2.39789527279837 -3.333      7.787  -0.428  0.67617
factor(log(dose + 1))3.52636052461616  3.333      7.787   0.428  0.67617
factor(log(dose + 1))4.61512051684126 21.000      7.787   2.697  0.01942 *
factor(log(dose + 1))5.8111409929767 15.667      7.787   2.012  0.06722 .
factor(log(dose + 1))6.90875477931522  8.000      7.787   1.027  0.32449
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 9.536 on 12 degrees of freedom
Multiple R-squared:  0.5475,    Adjusted R-squared:  0.359
F-statistic: 2.904 on 5 and 12 DF,  p-value: 0.06047
```

```
In [92]: anova(model_a, model_b)
```

| Res.Df | RSS      | Df | Sum of Sq | F        | Pr(>F)    |
|--------|----------|----|-----------|----------|-----------|
| 16     | 1881.063 | NA | NA        | NA       | NA        |
| 12     | 1091.333 | 4  | 789.7299  | 2.170913 | 0.1341968 |

```
In [93]: 1-pf(2.170913,4,12)
```

```
0.134196803857386
```

- Since p value is greater than 5% significance level, we accept null hypothesis. We can say that there is no lack of fit.

**Problem 2:** The `gammaray` data set from the `faraway` library includes X-ray decay light curve of Gamma ray burst 050525a obtained with the X-Ray Telescope (XRT) on board the Swift satellite. The data set has 63 brightness measurements in the 0.4-4.5 keV spectral band at times ranging from 2 minutes to 5 days after the burst. An appropriate model to predict flux as a function of time using appropriate weights is required.

(a) Find an appropriate linearizing transformation for either the response, the predictor or both if necessary.

(b) Fit a regression model to predict flux as a function of time, using the proposed transformation in part a) and appropriate weights. Note that the measurement error of the flux is available in the data set.

(c) Inspect the resulting model residuals and comment on your results

(d) In case the residuals do not have an appropriate behavior propose a new model for this data set.


```
In [95]: data(gammaray)
         head(gammaray,2)
```

| time | flux  | error |
|------|-------|-------|
| 133  | 122.7 | 5.7   |
| 143  | 109.5 | 5.4   |

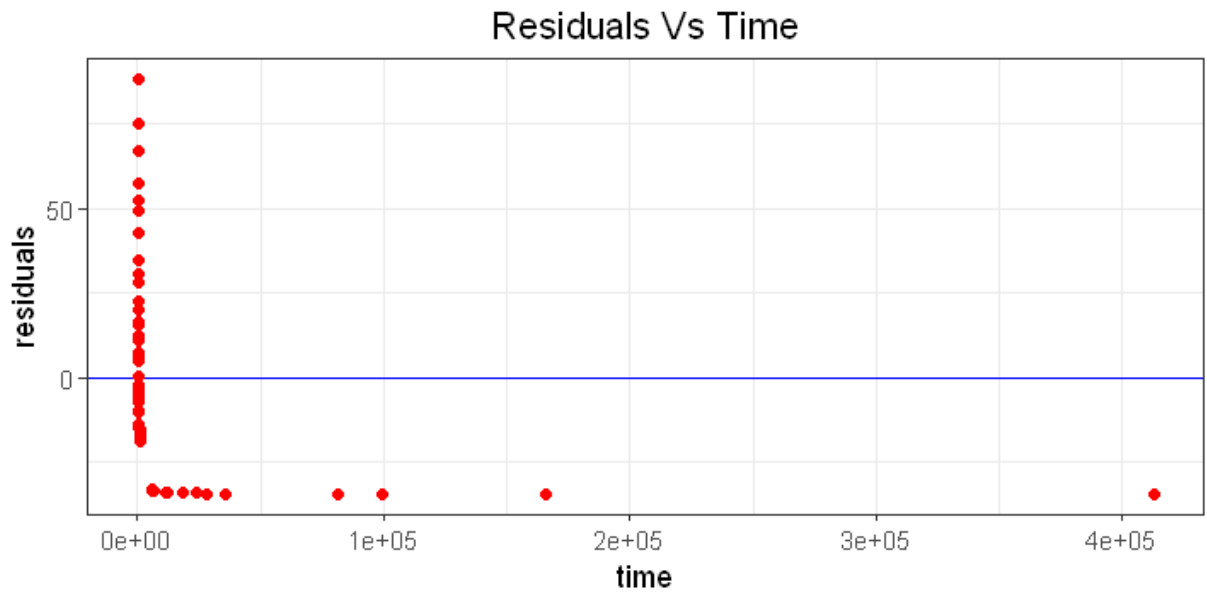
(a) Find an appropriate linearizing transformation for either the response, the predictor or both if necessary.

```
In [154]: mod_a = lm(flux ~ time, data = gammaray)
```

```
In [155]: ggplot(data = gammaray, aes(x = time, y = flux))+
  geom_point()+
  geom_point(color='red')+
  labs(title='Flux Vs Time', x='time', y = 'flux')+
  geom_abline(intercept = mod_a$coef[1], slope= mod_a$coef[2])
```




```
In [153]: ggplot(data = gammaray, aes(x = time, y = mod_a$res))+
  geom_point()+
  geom_hline(yintercept = 0, color = 'blue') +
  theme_update(plot.title = element_text(hjust = 0.5))+
  theme_set(theme_bw())+
  labs(title='Residuals Vs Time', x='time', y = 'residuals')+
  geom_point(color='red')
```




- Since the time is varying with exponentially, lets apply log to time

```
In [129]: mod_b = lm(flux ~ log(time), data = gammaray)
```


```
In [139]: ggplot(data = gammaray, aes(x = log(time), y = flux))+  
  geom_point()+  
  geom_point(color='red')+  
  labs(title='Flux Vs Time', x='log(time)', y = 'flux')+  
  geom_abline(intercept = mod_b$coef[1], slope= mod_b$coef[2])
```



```
In [137]: ggplot(data = gammaray, aes(x = log(time), y = mod_b$res))+  
  geom_point()+  
  geom_hline(yintercept = 0, color = 'blue') +  
  theme_update(plot.title = element_text(hjust = 0.5))+  
  theme_set(theme_bw())+  
  labs(title='Residuals Vs log(Time) mod_b', x='log(time)', y = 'residuals')+  
  geom_point(color='red')
```




```
In [213]: ggplot(data = gammaray, aes(x = mod_b$fitted, y = mod_b$res))+
  geom_point()+
  geom_hline(yintercept = 0, color = 'blue') +
  theme_update(plot.title = element_text(hjust = 0.5))+
  theme_set(theme_bw())+
  labs(title='Residuals Vs Fitted mod_b', x='fitted', y = 'residuals')+
  geom_point(color='red')
```




- The lm model is still not fitting well on the data, there appears to be log linear relationship.
- Lets also apply log transformation to response variable as well.

```
In [214]: mod_c = lm(log(flux) ~ log(time), data = gammaray)
```

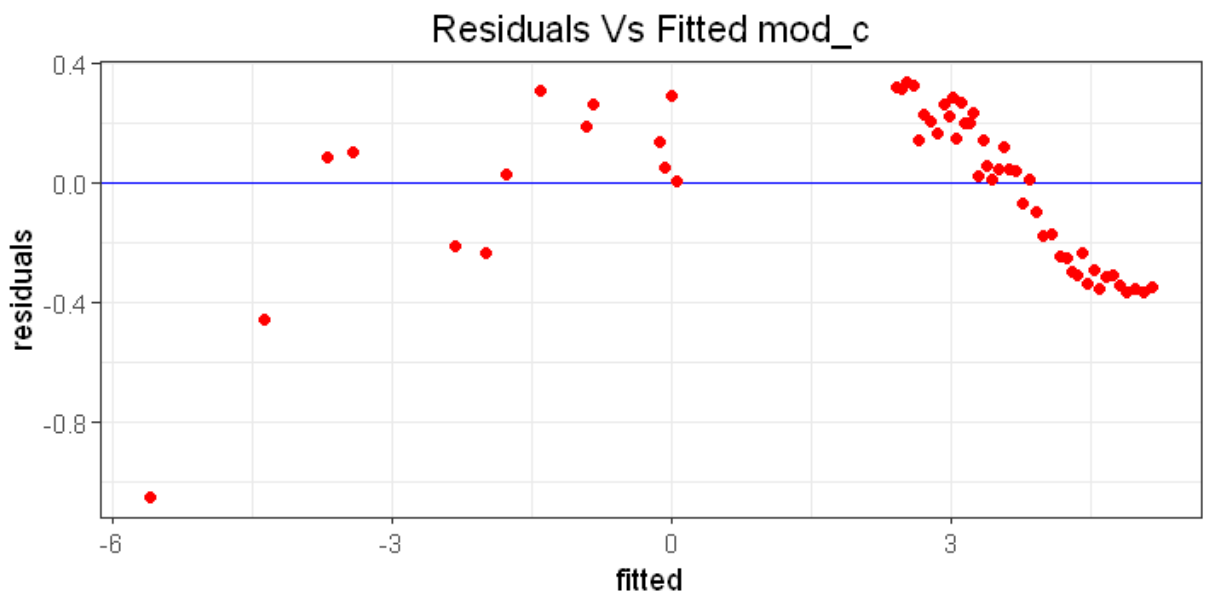
```
In [215]: ggplot(data = gammaray, aes(x = log(time), y = log(flux)))+  
  geom_point()+  
  geom_point(color='red')+  
  labs(title='Flux Vs Time', x='log(time)', y = 'log(flux)')+  
  geom_abline(intercept = mod_c$coef[1], slope= mod_c$coef[2])
```



```
In [216]: ggplot(data = gammaray, aes(x = log(time), y = mod_c$res))+
  geom_point()+
  geom_hline(yintercept = 0, color = 'blue') +
  theme_update(plot.title = element_text(hjust = 0.5))+
  theme_set(theme_bw())+
  labs(title='Residuals Vs log(Time) mod_c', x='log(time)', y = 'residuals')+
  geom_point(color='red')
```



```
In [219]: ggplot(data = gammaray, aes(x = mod_c$fitted, y = mod_c$res))+
  geom_point()+
  geom_hline(yintercept = 0, color = 'blue') +
  theme_update(plot.title = element_text(hjust = 0.5))+
  theme_set(theme_bw())+
  labs(title='Residuals Vs Fitted mod_c', x='fitted', y = 'residuals')+
  geom_point(color='red')
```



- We can see that transorming both response and predictor with log, the lm model is fitting better on the data.



- The residuals look like auto correlated with time. The value of a residual at a particular point depend upon value of preceding residual.

```
In [190]: dwtest(mod_c)
```

Durbin-Watson test

```
data: mod_c
DW = 0.2978, p-value < 2.2e-16
alternative hypothesis: true autocorrelation is greater than 0
```


- The low p value of Durbin-Watson test indicates that residuals are highly correlated and they depend upon value of preceding residual.

**(b) Fit a regression model to predict flux as a function of time, using the proposed transformation in part a) and appropriate weights. Note that the measurement error of the flux is available in the data set.**

```
In [140]: mod_d = lm(log(flux) ~ log(time), data = gammaray, weights = 1/error^2)
```


**(c) Inspect the resulting model residuals and comment on your results**

```
In [150]: ggplot(data = gammaray, aes(x = log(time), y = log(flux)))+
  geom_point()+
  geom_point(color='red')+
  labs(title='Flux Vs Time', x='log(time)', y = 'log(flux)')+
  geom_abline(intercept = mod_d$coef[1], slope= mod_d$coef[2], color='red')+
  geom_abline(intercept = mod_c$coef[1], slope= mod_c$coef[2], color='black')
```



- Red color line is the fitted line after using  $weights = \frac{1}{sd^2}$ .

```
In [192]: ggplot(data = gammaray, aes(x = log(time), y = mod_d$res))+
  geom_point()+
  geom_hline(yintercept = 0, color = 'blue') +
  theme_update(plot.title = element_text(hjust = 0.5))+
  theme_set(theme_bw())+
  labs(title='Residuals Vs log(Time) mod_d', x='log(time)', y = 'residuals')+
  geom_point(color='red')
```



- When we use the weights (variance) of  $X_i$ , the lm model is fitted with below condition.

The GLS estimate of  $\beta$  minimizes:

$$(\mathbf{y} - \mathbf{X}\beta)^\top \Sigma^{-1} (\mathbf{y} - \mathbf{X}\beta) = \sum_{i=1}^n \frac{(y_i - \mathbf{x}_i^\top \beta)^2}{\sigma_i^2}$$

- Since we are using  $weights = \frac{1}{sd^2}$ , More the variance at a particular  $X_i$ , lesser weightage will be given to that point while fitting the model.
- We can see that the residuals are still highly correlated. The value of a residual at a particular point depend upon value of preceding residual.

**(d) In case the residuals do not have an appropriate behavior propose a new model for this data set.**

- We can fit a regression model with autocorrelated errors.

```
In [185]: library(nlme)
mod_e = gls(log(flux) ~ log(time), correlation = corARMA(p=1), data = gammaray)
summary(mod_e)
```

Generalized least squares fit by REML

Model: log(flux) ~ log(time)

Data: gammaray

|  | AIC       | BIC       | logLik   |
|--|-----------|-----------|----------|
|  | -56.17774 | -47.73424 | 32.08887 |

Correlation Structure: AR(1)

Formula: ~1

Parameter estimate(s):

Phi  
0.9915039

Coefficients:

|             | Value     | Std.Error | t-value   | p-value |
|-------------|-----------|-----------|-----------|---------|
| (Intercept) | 12.735637 | 1.0804918 | 11.78689  | 0       |
| log(time)   | -1.524327 | 0.0580027 | -26.28027 | 0       |

Correlation:

(Intr)  
log(time) -0.456

Standardized residuals:

|  | Min        | Q1         | Med       | Q3        | Max       |
|--|------------|------------|-----------|-----------|-----------|
|  | -0.4390054 | -0.1295899 | 0.3330974 | 0.5231954 | 1.1090041 |

Residual standard error: 1.081339

Degrees of freedom: 63 total; 61 residual

```
In [186]: intervals(mod_e)
```

Approximate 95% confidence intervals

Coefficients:

|             | lower     | est.      | upper     |
|-------------|-----------|-----------|-----------|
| (Intercept) | 10.575060 | 12.735637 | 14.896214 |
| log(time)   | -1.640311 | -1.524327 | -1.408344 |

attr(,"label")  
[1] "Coefficients:"

Correlation structure:

|     | lower      | est.      | upper     |
|-----|------------|-----------|-----------|
| Phi | -0.9605414 | 0.9915039 | 0.9999993 |

attr(,"label")  
[1] "Correlation structure:"

Residual standard error:

|  | lower      | est.       | upper        |
|--|------------|------------|--------------|
|  | 0.01021064 | 1.08133890 | 114.51720689 |

**Problem 3: The divusa data set from the faraway library reports the divorce rate in the USA**

from 1920-1996. Fit a regression model with divorce as the response and unemployed, femlab, marriage, birth and military as predictors.

- (a) Make a residual plot against year. Do you think the residuals are auto-correlated? Why?
- (b) Make a formal test for the autocorrelation of the residuals.
- (c) Assume the residuals behave as an autoregressive model of order 1 (AR(1)). Use the function `gls` from the library `nlme` to fit a regression model to the variable `divorce` as a response and the `unemployed`, `femlab`, `marriage`, `birth` and `military` as predictors. Use the option `correlation = corAR1( form = ~ year)` and `method = "ML"`. Make sure you do `>help(gls)` before fitting the model, in order to understand the meaning of the parameters `correlation` and `method` in the `gls` function call.
- (d) What is the estimated autocorrelation coefficient? Is this coefficient significant? Hint: Use function `intervals` to get the CI.
- (e) Discuss whether the GLS model change which variables are found significant when compared to the LM model.
- (f) Comment on why there might be correlation in the errors.

In [223]:

```
data(divusa)
head(divusa,2)
```

| year | divorce | unemployed | femlab | marriage | birth | military |
|------|---------|------------|--------|----------|-------|----------|
| 1920 | 8.0     | 5.2        | 22.70  | 92       | 117.9 | 3.2247   |
| 1921 | 7.2     | 11.7       | 22.79  | 83       | 119.8 | 3.5614   |

```
In [242]: mod_a = lm(divorce ~ unemployed + femlab + marriage + birth + military, data = di
summary(mod_a)
```

Call:

```
lm(formula = divorce ~ unemployed + femlab + marriage + birth +
    military, data = divusa)
```

Residuals:

|  | Min     | 1Q      | Median  | 3Q     | Max    |
|--|---------|---------|---------|--------|--------|
|  | -3.8611 | -0.8916 | -0.0496 | 0.8650 | 3.8300 |

Coefficients:

|             | Estimate | Std. Error | t value | Pr(> t )     |
|-------------|----------|------------|---------|--------------|
| (Intercept) | 2.48784  | 3.39378    | 0.733   | 0.4659       |
| unemployed  | -0.11125 | 0.05592    | -1.989  | 0.0505 .     |
| femlab      | 0.38365  | 0.03059    | 12.543  | < 2e-16 ***  |
| marriage    | 0.11867  | 0.02441    | 4.861   | 6.77e-06 *** |
| birth       | -0.12996 | 0.01560    | -8.333  | 4.03e-12 *** |
| military    | -0.02673 | 0.01425    | -1.876  | 0.0647 .     |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 1.65 on 71 degrees of freedom

Multiple R-squared: 0.9208, Adjusted R-squared: 0.9152


F-statistic: 165.1 on 5 and 71 DF, p-value: < 2.2e-16

**(a) Make a residual plot against year. Do you think the residuals are auto-correlated? Why?**

```
In [231]: ggplot(data = divusa, aes(x = divusa$year, y = mod_a$res))+
  geom_point()+
  geom_hline(yintercept = 0, color = 'blue') +
  theme_update(plot.title = element_text(hjust = 0.5))+
  theme_set(theme_bw())+
  labs(title='Residuals Vs year', x='year', y = 'residuals')+
  geom_point(color='red')+
  geom_line()
```



```
In [225]: ggplot(data = divusa, aes(x = mod_a$fitted, y = mod_a$res))+
  geom_point()+
  geom_hline(yintercept = 0, color = 'blue') +
  theme_update(plot.title = element_text(hjust = 0.5))+
  theme_set(theme_bw())+
  labs(title='Residuals Vs fitted mod_a', x='fitted', y = 'residuals')+
  geom_point(color='red')+
  geom_line()
```



- We can see that the residuals are still highly correlated. They follow a trend. The value of a residual at a particular point depends upon value of preceding residual.

**(b) Make a formal test for the autocorrelation of the residuals.**

In [226]: `dwtest(mod_a)`

Durbin-Watson test

```
data: mod_a
DW = 0.29988, p-value < 2.2e-16
alternative hypothesis: true autocorrelation is greater than 0
```

- Here, the null hypothesis is that the errors are not correlated, and alternative hypothesis is that the errors are correlated.
- By performing Durbin-Watson test, we got a very low p value. Hence we reject the null hypothesis and accept that the errors are highly correlated.

**(c) Assume the residuals behave as an autoregressive model of order 1 (AR(1)). Use the function `gls` from the library `nlme` to fit a regression model to the variable `divorce` as a response and the `unemployed`, `femlab`, `marriage`, `birth` and `military` as predictors. Use the option `correlation = corAR1( form = ~ year)` and `method = "ML"`. Make sure you do `>help(gls)` before fitting the model, in order to understand the meaning of the parameters `correlation` and `method` in the `gls` function call.**

```
In [233]: mod_b = gls(divorce ~ unemployed + femlab + marriage + birth + military, correlation = 0.9715486,
                    method = "ML", data = divusa)
summary(mod_b)
```

Generalized least squares fit by maximum likelihood

Model: divorce ~ unemployed + femlab + marriage + birth + military

Data: divusa

|  | AIC      | BIC      | logLik    |
|--|----------|----------|-----------|
|  | 179.9523 | 198.7027 | -81.97613 |

Correlation Structure: AR(1)

Formula: ~year

Parameter estimate(s):

Phi  
0.9715486

Coefficients:

|             | Value     | Std.Error | t-value   | p-value |
|-------------|-----------|-----------|-----------|---------|
| (Intercept) | -7.059682 | 5.547193  | -1.272658 | 0.2073  |
| unemployed  | 0.107643  | 0.045915  | 2.344395  | 0.0219  |
| femlab      | 0.312085  | 0.095151  | 3.279878  | 0.0016  |
| marriage    | 0.164326  | 0.022897  | 7.176766  | 0.0000  |
| birth       | -0.049909 | 0.022012  | -2.267345 | 0.0264  |
| military    | 0.017946  | 0.014271  | 1.257544  | 0.2127  |

Correlation:

|            | (Intr) | unmply | femlab | marrig | birth |
|------------|--------|--------|--------|--------|-------|
| unemployed | -0.420 |        |        |        |       |
| femlab     | -0.802 | 0.240  |        |        |       |
| marriage   | -0.516 | 0.607  | 0.307  |        |       |
| birth      | -0.379 | 0.041  | 0.066  | -0.094 |       |
| military   | -0.036 | 0.436  | -0.311 | 0.530  | 0.128 |

Standardized residuals:

|  | Min        | Q1         | Med        | Q3        | Max       |
|--|------------|------------|------------|-----------|-----------|
|  | -1.4509327 | -0.9760939 | -0.6164694 | 1.1375377 | 2.1593261 |

Residual standard error: 2.907664

Degrees of freedom: 77 total; 71 residual

**(d) What is the estimated autocorrelation coefficient? Is this coefficient significant? Hint: Use function intervals to get the CI.**

- We know that residuals follow a model of the form:  $e_{t+1} = \phi e_t + \gamma_t$ . Here,  $\phi$  is the autocorrelation coefficient.
- By fitting generalized least squares model with correlation, we got  $\phi$  value as 0.9715486.
- To check its significance, we can check confidence interval of  $\phi$ .



```
In [234]: intervals(mod_b)
```

Approximate 95% confidence intervals

```
Coefficients:
              lower      est.      upper
(Intercept) -18.12047043 -7.05968166  4.001107109
unemployed   0.01609101  0.10764313  0.199195251
femlab       0.12235846  0.31208493  0.501811412
marriage     0.11867101  0.16432630  0.209981587
birth       -0.09380023 -0.04990919 -0.006018159
military    -0.01050915  0.01794640  0.046401944
attr(,"label")
[1] "Coefficients:"
```

```
Correlation structure:
              lower      est.      upper
Phi 0.6529393 0.9715486 0.9980183
attr(,"label")
[1] "Correlation structure:"
```

```
Residual standard error:
              lower      est.      upper
0.7976096  2.9076645 10.5998135
```

- We can see that the 95% CI of  $\phi$  is greater than 0. Hence this coefficient is significant.

**(e) Discuss whether the GLS model change which variables are found significant when compared to the LM model.**

```
In [240]: summary(mod_a)
```

Call:

```
lm(formula = divorce ~ unemployed + femlab + marriage + birth +  
    military, data = divusa)
```

Residuals:

|  | Min     | 1Q      | Median  | 3Q     | Max    |
|--|---------|---------|---------|--------|--------|
|  | -3.8611 | -0.8916 | -0.0496 | 0.8650 | 3.8300 |

Coefficients:

|             | Estimate | Std. Error | t value | Pr(> t )     |
|-------------|----------|------------|---------|--------------|
| (Intercept) | 2.48784  | 3.39378    | 0.733   | 0.4659       |
| unemployed  | -0.11125 | 0.05592    | -1.989  | 0.0505 .     |
| femlab      | 0.38365  | 0.03059    | 12.543  | < 2e-16 ***  |
| marriage    | 0.11867  | 0.02441    | 4.861   | 6.77e-06 *** |
| birth       | -0.12996 | 0.01560    | -8.333  | 4.03e-12 *** |
| military    | -0.02673 | 0.01425    | -1.876  | 0.0647 .     |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.65 on 71 degrees of freedom

Multiple R-squared: 0.9208, Adjusted R-squared: 0.9152

F-statistic: 165.1 on 5 and 71 DF, p-value: < 2.2e-16

- In lm model, femlab, marriage, and birth are found to be significant at 5% significance level.

In [243]: `summary(mod_b)`

```
Generalized least squares fit by maximum likelihood
Model: divorce ~ unemployed + femlab + marriage + birth + military
Data: divusa
      AIC      BIC    logLik
179.9523 198.7027 -81.97613
```

```
Correlation Structure: AR(1)
Formula: ~year
Parameter estimate(s):
  Phi
0.9715486
```

```
Coefficients:
      Value Std.Error  t-value p-value
(Intercept) -7.059682  5.547193 -1.272658  0.2073
unemployed   0.107643  0.045915  2.344395  0.0219
femlab       0.312085  0.095151  3.279878  0.0016
marriage     0.164326  0.022897  7.176766  0.0000
birth       -0.049909  0.022012 -2.267345  0.0264
military     0.017946  0.014271  1.257544  0.2127
```

```
Correlation:
      (Intr) unmply femlab marrig birth
unemployed -0.420
femlab      -0.802  0.240
marriage    -0.516  0.607  0.307
birth       -0.379  0.041  0.066 -0.094
military    -0.036  0.436 -0.311  0.530  0.128
```

```
Standardized residuals:
      Min      Q1      Med      Q3      Max
-1.4509327 -0.9760939 -0.6164694  1.1375377  2.1593261
```

```
Residual standard error: 2.907664
Degrees of freedom: 77 total; 71 residual
```

- With GLS model, unemployed, femlab, marriage and birth are found to be significant at 5% significance level.
- Compared to linear model, unemployed is also found to be significant in generalized linear sqaures model.

**(f) Comment on why there might be correlation in the errors.**

- There might be correlation in the errors due to following reason:
  1. The correlation between errors occurs mostly in time series data. The observation at a particular time  $t$  is affected by observation taken at time  $t - 1$  or  $t - 2$ . This observations affects the residuals and hence they become correlated.

