

Problem 1: Take the fat data and use the percentage of body fat, siri, as the response, and the other variables except brozek and density as potential predictors. Remove every tenth observation from the data for use as a test sample. Use the remaining data as a training sample building the following models:

(a) Linear Regression with all predictors

(b) Linear regression with variables selected using BIC by using stepwise elimination.

(c) Principal Components Regression by selecting the number of PC's using Cross-Validation.


(d) Lasso regression

(e) Use the previous models you fitted in (a) to (d) to predict the response in the test sample. Make a report on the models' performance by using the Root Mean Square Error (RMSE). Which model performed best?

```
In [2]: library('faraway')
```

```
In [6]: data(fat)
        head(fat,2)
```

brozek	siri	density	age	weight	height	adipos	free	neck	chest	abdom	hip	thigh	knee
12.6	12.3	1.0708	23	154.25	67.75	23.7	134.9	36.2	93.1	85.2	94.5	59.0	37.3
6.9	6.1	1.0853	22	173.25	72.25	23.4	161.3	38.5	93.6	83.0	98.7	58.7	37.3



```
In [12]: test_data_idx = seq(from = 10, to = nrow(fat), by = 10)
        fat_trn = fat[-test_data_idx,]
        fat_tst = fat[test_data_idx,]
```

(a) Build Linear Regression with all predictors

```
In [15]: mod_lin = lm(siri ~ .-brozek-density, fat_trn)
```

```
In [16]: summary(mod_lin)
```

Call:

```
lm(formula = siri ~ . - brozek - density, data = fat_trn)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-5.8314	-0.6722	0.1828	0.9150	6.6619

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-12.591885	6.448868	-1.953	0.052193	.
age	0.007978	0.012320	0.648	0.517983	
weight	0.362999	0.023314	15.570	< 2e-16	***
height	0.049026	0.040315	1.216	0.225315	
adipos	-0.514032	0.114074	-4.506	1.09e-05	***
free	-0.564773	0.014889	-37.933	< 2e-16	***
neck	0.016525	0.089863	0.184	0.854272	
chest	0.120219	0.039590	3.037	0.002694	**
abdom	0.140108	0.042186	3.321	0.001056	**
hip	0.006197	0.056101	0.110	0.912148	
thigh	0.195057	0.054460	3.582	0.000424	***
knee	0.106637	0.093534	1.140	0.255542	
ankle	0.125118	0.081303	1.539	0.125325	
biceps	0.096199	0.064656	1.488	0.138278	
forearm	0.230775	0.073332	3.147	0.001888	**
wrist	0.139279	0.206804	0.673	0.501378	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.55 on 211 degrees of freedom

Multiple R-squared: 0.9692, Adjusted R-squared: 0.967

F-statistic: 442.5 on 15 and 211 DF, p-value: < 2.2e-16

(b) Build Linear regression with variables selected using BIC by using stepwise elimination.

```
In [26]: # install.packages('leaps')
library(leaps)
mod_lin_vs = regsubsets(siri ~ .-brozek-density, fat_trn)
mod_lin_vs_summary = summary(mod_lin_vs)
mod_lin_vs_summary$which
```

(Intercept)	age	weight	height	adipos	free	neck	chest	abdom	hip	thigh	knee
TRUE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE	FALSE	FALSE	FALSE
TRUE	FALSE	TRUE	FALSE	FALSE	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
TRUE	FALSE	TRUE	FALSE	FALSE	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
TRUE	FALSE	TRUE	FALSE	FALSE	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE
TRUE	FALSE	TRUE	FALSE	TRUE	TRUE	FALSE	FALSE	TRUE	FALSE	FALSE	FALSE
TRUE	FALSE	TRUE	FALSE	TRUE	TRUE	FALSE	FALSE	TRUE	FALSE	TRUE	FALSE
TRUE	FALSE	TRUE	FALSE	TRUE	TRUE	FALSE	TRUE	TRUE	FALSE	TRUE	FALSE
TRUE	FALSE	TRUE	FALSE	TRUE	TRUE	FALSE	TRUE	TRUE	FALSE	TRUE	FALSE



```
In [57]: n = nrow(fat_trn)
msize = 2:9
BICs = n*log(mod_lin_vs_summary$rss/n) + msize*log(n)
BICs
which.min(BICs)
```


```
734.955027999442 272.966016385392 264.5296194184 261.250401883378
249.641209684591 243.26369250278 238.714817402774 240.045098659401
```

```

In [135]: library('ggplot2')
options(repr.plot.width=6, repr.plot.height=3)

plot1 <- ggplot(data = data.frame(msize, BICs), aes(x = msize-1, y = BICs, size = 
  geom_point()+
  theme_update(plot.title = element_text(hjust = 0.5))+
  theme_set(theme_bw()))+
  geom_point(color = 'dodgerblue')+
  labs(title='Variable Selection by BIC', x='No. of parameters', y = 'BIC')
plot1

```



```

In [37]: mod_final_cv = lm(siri ~ weight+adipos+free+chest+abdom+thigh+forearm, fat_trn)

```

```
In [41]: summary(mod_final_cv)
```

Call:

```
lm(formula = siri ~ weight + adipos + free + chest + abdom +  
    thigh + forearm, data = fat_trn)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-6.0226	-0.5667	0.2136	0.8846	7.3072

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.10003	3.27933	-0.335	0.737612
weight	0.38940	0.01904	20.449	< 2e-16 ***
adipos	-0.56961	0.09277	-6.140	3.84e-09 ***
free	-0.55874	0.01418	-39.410	< 2e-16 ***
chest	0.12036	0.03837	3.136	0.001945 **
abdom	0.14634	0.03872	3.779	0.000203 ***
thigh	0.19170	0.04358	4.399	1.69e-05 ***
forearm	0.28758	0.06689	4.299	2.58e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.565 on 219 degrees of freedom

Multiple R-squared: 0.9674, Adjusted R-squared: 0.9663

F-statistic: 927.8 on 7 and 219 DF, p-value: < 2.2e-16

(c) Principal Components Regression by selecting the number of PC's using Cross-Validation.

```
In [405]: # install.packages('pls')  
library(pls)  
mod_pcr = pcr(siri ~ .-brozek-density, data = fat_trn, ncomp=15, scale = TRUE)  
summary(mod_pcr)
```

Data: X dimension: 227 15

Y dimension: 227 1

Fit method: svdpc

Number of components considered: 15

TRAINING: % variance explained


	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
X	62.90	73.17	80.41	84.99	88.98	91.43	93.60	95.42
siri	36.65	62.19	63.67	63.71	65.59	82.69	84.05	84.06
	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	
X	96.74	97.95	98.80	99.31	99.62	99.88	100.00	
siri	91.02	91.86	92.52	92.77	93.76	93.81	96.92	

Select optimum number of components

```
In [407]: fat_pca <- prcomp(fat_trn[, -3:-1], scale = TRUE)

options(repr.plot.width=6, repr.plot.height=3)

plot1 <- ggplot(data = data.frame(1:15, fat_pca$sdev[1:15]), aes(x = 1:15, y = fat_pca$sdev)) +
  geom_line() +
  theme_update(plot.title = element_text(hjust = 0.5)) +
  theme_set(theme_bw()) +
  geom_point(color = 'dodgerblue') +
  labs(title='# of PC vs PCAs std dev', x='PCA number', y = 'PCAs Std Dev')
plot1
```



- From the above plot, we can see that after 4 number of PCAs, there is significant decrease in cumulative % variation explained.


Selection PCA which minimizes RMSE


```
In [469]: #Use 10-fold Cross-Validation
set.seed(123)
mod_pcr_cv = pcr(siri ~ .-brozek-density, data=fat_trn, validation="CV", ncomp=15,

pcrCV = RMSEP(mod_pcr_cv, estimate="CV")

options(repr.plot.width=6, repr.plot.height=4)

plot(pcrCV)
validationplot(mod_pcr_cv, val.type = "R2")
axis(side=1, at=seq(0, 15, by=1))
```





```
In [442]: which.min(pcrCV$val)
          pcrCV$val
```

14

```
8.5508878521787  6.86395311956274  5.34194412745452  5.34797228210977
5.38801861903322  5.51830445765396  4.29674505329343  3.76131767286886
3.51903670424762  3.47616501717872  3.40451896668907  3.28359862327293
3.27474528240373  3.24783052900413  4.07177003151727  3.59971517373094
```

- With 14 principle components, we get min RMSE.
- But, 10 principle components explain more than 80% variation.

(d) Build Lasso regression

```
In [176]: fat_trn_x = as.matrix(fat_trn[, -3:-1])
          fat_tst_x = as.matrix(fat_tst[, -3:-1])
```

- Build cross validated lasso regression model


```
In [187]: library('glmnet')
mod_cv_lasso = cv.glmnet(x=fat_trn_x , y=fat_trn$siri, alpha = 1)
coef(mod_cv_lasso, s=mod_cv_lasso$lambda.min)
```


16 x 1 sparse Matrix of class "dgCMatrix"

```

              1
(Intercept) -12.77985378
age          .
weight       0.29716332
height       0.09248681
adipos       .
free         -0.50465875
neck         .
chest        0.04253080
abdom        0.18228390
hip          .
thigh        0.10130032
knee         0.20691171
ankle        0.04657393
biceps       0.07585771
forearm      0.21277081
wrist        .
```

```
In [196]: options(repr.plot.width=7, repr.plot.height=5)

plot(mod_cv_lasso)
```



```
In [202]: mod_cv_lasso$lambda.min  
log(mod_cv_lasso$lambda.min)
```

0.0875960720583751

-2.43501912156138

- Here value of lambda that achieves min. mean squared error is 0.0875960720583751.

(e) Use the previous models you fitted in (a) to (d) to predict the response in the test sample. Make a report on the models' performance by using the Root Mean Square Error (RMSE). Which model performed best?

- Write function to calculate RMSE

```
In [144]: calc_RMSE = function(pred, actual){  
  sqrt(mean((pred - actual)**2))  
}
```

RMSE (a)

```
In [150]: pred = predict(mod_lin, fat_tst)  
calc_RMSE(pred, fat_tst$siri)
```

1.13152878296094

RMSE (b)

```
In [151]: pred = predict(mod_final_cv, fat_tst)  
calc_RMSE(pred, fat_tst$siri)
```

1.23391847665925

RMSE (c)

```
In [479]: pred = predict(mod_pcr_cv, fat_tst, ncomp=10)  
calc_RMSE(pred, fat_tst$siri)
```

1.77337949292342

RMSE (d)

```
In [203]: pred = predict(mod_cv_lasso, fat_tst_x, s=mod_cv_lasso$lambda.min)
          calc_RMSE(pred, fat_tst$siri)
```

1.0991337382338

- As we can see, model with cross validated PCA performed best giving RMSE of 1.0488508. The cross validated lasso regression also performed better than linear models.

Problem 2: Use the chickwts data to fit a one-way ANOVA with weights as the response and feed as the predictor.

(a) Determine whether there are any differences in the weights of chickens according to their feed.

(b) Perform all necessary model diagnostics.

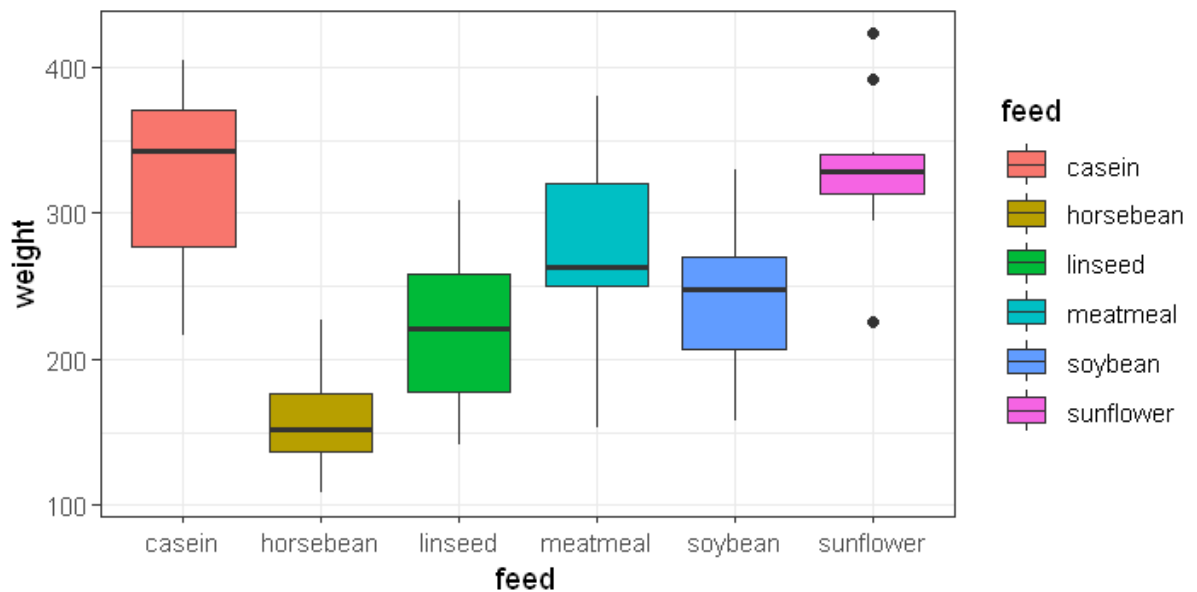
```
In [216]: data(chickwts)
          head(chickwts, 3)
          levels(chickwts$feed)
```

weight	feed
179	horsebean
160	horsebean
136	horsebean

'casein' 'horsebean' 'linseed' 'meatmeal' 'soybean' 'sunflower'

(a) Determine whether there are any differences in the weights of chickens according to their feed.

```
In [217]: options(repr.plot.width=6, repr.plot.height=3)
ggplot(chickwts, aes(x = feed, y = weight, fill=feed)) +
  geom_boxplot()
```



```
In [220]: summary(chickwts)
```

weight	feed
Min. :108.0	casein :12
1st Qu.:204.5	horsebean:10
Median :258.0	linseed :12
Mean :261.3	meatmeal :11
3rd Qu.:323.5	soybean :14
Max. :423.0	sunflower:12

```
In [465]: feeds = levels(chickwts$feed)
calc_mean = function(feed){
  mean(chickwts$weight[chickwts$feed==feed])
}

means = sapply(feeds, calc_mean)
means
```

casein	323.583333333333
horsebean	160.2
linseed	218.75
meatmeal	276.909090909091
soybean	246.428571428571
sunflower	328.916666666667

```
In [481]: mod = lm(weight ~ feed-1, data = chickwts)
summary(mod)
anova(mod)
```

Call:

```
lm(formula = weight ~ feed - 1, data = chickwts)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-123.909	-34.413	1.571	38.170	103.091

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
feedcasein	323.58	15.83	20.436	< 2e-16 ***
feedhorsebean	160.20	17.35	9.236	1.91e-13 ***
feedlinseed	218.75	15.83	13.815	< 2e-16 ***
feedmeatmeal	276.91	16.54	16.744	< 2e-16 ***
feedsoybean	246.43	14.66	16.810	< 2e-16 ***
feedsunflower	328.92	15.83	20.773	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 54.85 on 65 degrees of freedom

Multiple R-squared: 0.9629, Adjusted R-squared: 0.9595

F-statistic: 281.4 on 6 and 65 DF, p-value: < 2.2e-16

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
feed	6	5079211	846535.163	281.3761	1.677742e-44
Residuals	65	195556	3008.554	NA	NA

- Since the p value is lower than 0.05, we conclude that weights are different for different feeds.

```
In [235]: R_squared = var(mod$fitted.values)/var(chickwts$weight)
R_squared
```

0.541685465673916

```
In [236]: confint(mod)
```

	2.5 %	97.5 %
feedcasein	291.9608	355.2058
feedhorsebean	125.5593	194.8407
feedlinseed	187.1275	250.3725
feedmeatmeal	243.8805	309.9377
feedsoybean	217.1518	275.7053
feedsunflower	297.2942	360.5392

```
In [237]: # Pairwise t-test with Bonferroni Correction
pairwise.t.test(chickwts$weight,chickwts$feed,p.adjust.method = "bonferroni")
```

Pairwise comparisons using t tests with pooled SD

data: chickwts\$weight and chickwts\$feed

	casein	horsebean	linseed	meatmeal	soybean
horsebean	3.1e-08	-	-	-	-
linseed	0.00022	0.22833	-	-	-
meatmeal	0.68350	0.00011	0.20218	-	-
soybean	0.00998	0.00487	1.00000	1.00000	-
sunflower	1.00000	1.2e-08	9.3e-05	0.39653	0.00447

P value adjustment method: bonferroni

- Here, we can see that except linseed-horsebean, meatmeal-casein, linseed-meatmeal, linseed-soybean, meatmeal-soybean, casein-sunflower, meatmeal-sunflower, rest all differ significantly.
- We can also verify the same using Tukey comparisons

```
In [233]: # Compute the analysis of variance
res.aov <- aov(weight ~ feed-1, data = chickwts)
# Summary of the analysis
summary(res.aov)
TukeyHSD(res.aov)
```

```
              Df Sum Sq Mean Sq F value Pr(>F)
feed           6 5079211  846535   281.4 <2e-16 ***
Residuals    65  195556    3009
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


```
Tukey multiple comparisons of means
 95% family-wise confidence level
```

```
Fit: aov(formula = weight ~ feed - 1, data = chickwts)
```

```
$feed
```


	diff	lwr	upr	p adj
horsebean-casein	-163.383333	-232.346876	-94.41979	0.0000000
linseed-casein	-104.833333	-170.587491	-39.07918	0.0002100
meatmeal-casein	-46.674242	-113.906207	20.55772	0.3324584
soybean-casein	-77.154762	-140.517054	-13.79247	0.0083653
sunflower-casein	5.333333	-60.420825	71.08749	0.9998902
linseed-horsebean	58.550000	-10.413543	127.51354	0.1413329
meatmeal-horsebean	116.709091	46.335105	187.08308	0.0001062
soybean-horsebean	86.228571	19.541684	152.91546	0.0042167
sunflower-horsebean	168.716667	99.753124	237.68021	0.0000000
meatmeal-linseed	58.159091	-9.072873	125.39106	0.1276965
soybean-linseed	27.678571	-35.683721	91.04086	0.7932853
sunflower-linseed	110.166667	44.412509	175.92082	0.0000884
soybean-meatmeal	-30.480519	-95.375109	34.41407	0.7391356
sunflower-meatmeal	52.007576	-15.224388	119.23954	0.2206962
sunflower-soybean	82.488095	19.125803	145.85039	0.0038845

```
In [279]: options(repr.plot.width=6, repr.plot.height=5)
plot(TukeyHSD(res.aov), las=1, cex.axis = 0.5, col = 'blue')
```



(b) Perform all necessary model diagnostics.

```
In [282]: options(repr.plot.width=4, repr.plot.height=4)
plot(mod, 1)
```




```
In [269]: # install.packages('car')
library(car)
leveneTest(weight ~ feed - 1, data = chickwts, center=mean)
```

	Df	F value	Pr(>F)
group	5	0.987329	0.4324101
	65	NA	NA

- In the residuals vs fitted plot, the errors seem to have zero mean and constant variances. Also, from leveneTest, the p value is greater than 5% α level, hence we can accept null hypothesis.

```
In [257]: plot(mod, 2)
```




```
In [264]: # Shapiro-Wilk test  
shapiro.test(x = mod$residuals)
```

Shapiro-Wilk normality test


data: mod\$residuals
W = 0.98616, p-value = 0.6272

- From the plot, as well as Shapiro-Wilk test, we can assume normality of the residuals.

```
In [283]: plot(mod,3)
```



```
In [284]: plot(mod, 4)
```



- The maximum cook's distance is ~0.09 which is within limits. Hence, we conclude that there are no outliers.

Problem 3: Use the infmort data to fit a one-way ANOVA with income as the response and

region as the predictor.

(a) Determine whether income varies with region. Perform all necessary model diagnostics.

(b) In case income varies with region, determine which pairs of regions are different.

(c) In case you need a transformation for the response, re-fit the model and make a comparison with the previous results.


```
In [291]: data(infmort)
          head(infmort,2)
          levels(infmort$region)
```

	region	income	mortality	oil
Australia	Asia	3426	26.7	no oil exports
Austria	Europe	3350	23.7	no oil exports

'Africa' 'Europe' 'Asia' 'Americas'

(a) Determine whether income varies with region. Perform all necessary model diagnostics.

```
In [293]: options(repr.plot.width=6, repr.plot.height=3)
          ggplot(infmort, aes(x = region, y = income, fill=region)) +
            geom_boxplot()
```



```
In [496]: mod = lm(income ~ region-1, data = infmort)
summary(mod)
anova(mod)
```

Call:
lm(formula = income ~ region - 1, data = infmort)

Residuals:

Min	1Q	Median	3Q	Max
-2634.2	-515.9	-192.2	7.8	4583.1

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
regionAfrica	273.2	180.5	1.514	0.13317
regionEurope	3040.2	248.0	12.257	< 2e-16 ***
regionAsia	638.9	192.1	3.325	0.00123 **
regionAmericas	939.9	219.4	4.283	4.21e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1052 on 101 degrees of freedom
Multiple R-squared: 0.643, Adjusted R-squared: 0.6289
F-statistic: 45.48 on 4 and 101 DF, p-value: < 2.2e-16

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
region	4	201473156	50368289	45.47927	8.592504e-22
Residuals	101	111857493	1107500	NA	NA

```
In [497]: R_squared = var(mod$fitted.values)/var(infmort$income)
R_squared
```

0.464120440014417

```
In [498]: confint(mod)
```

	2.5 %	97.5 %
regionAfrica	-84.79121	631.2618
regionEurope	2548.16186	3532.2826
regionAsia	257.71835	1020.0150
regionAmericas	504.56751	1375.1716

```
In [488]: # Pairwise t-test with Bonferroni Correction
pairwise.t.test(infmort$income, infmort$region, p.adjust.method = "bonferroni")
```

Pairwise comparisons using t tests with pooled SD

data: infmort\$income and infmort\$region

	Africa	Europe	Asia
Europe	7.7e-14	-	-
Asia	1.00	7.2e-11	-
Americas	0.13	3.9e-08	1.00

P value adjustment method: bonferroni

```
In [494]: # Compute the analysis of variance
res.aov <- aov(income ~ region-1, data = infmort)
# Summary of the analysis
summary(res.aov)
TukeyHSD(res.aov)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
region	4	201473156	50368289	45.48	<2e-16 ***
Residuals	101	111857493	1107500		


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Tukey multiple comparisons of means
95% family-wise confidence level

Fit: aov(formula = income ~ region - 1, data = infmort)

\$region		diff	lwr	upr	p adj
Europe-Africa		2766.9869	1965.63247	3568.341	0.0000000
Asia-Africa		365.6314	-323.00285	1054.266	0.5103789
Americas-Africa		666.6343	-75.58589	1408.854	0.0945000
Asia-Europe		-2401.3556	-3220.99422	-1581.717	0.0000000
Americas-Europe		-2100.3527	-2965.50075	-1235.205	0.0000000
Americas-Asia		301.0029	-460.92186	1062.928	0.7311095

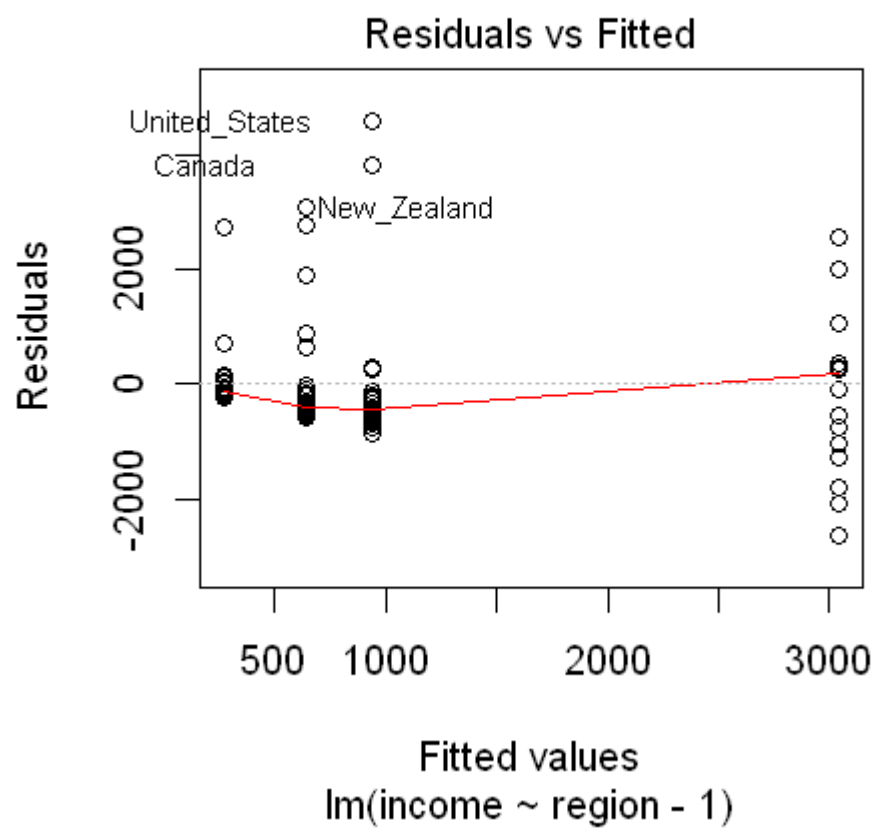
```
In [493]: options(repr.plot.width=6, repr.plot.height=3)
plot(TukeyHSD(res.aov), las=1, cex.axis = 0.5, col = 'blue')
```



Here, we can see that Africa-Europe, Europe-Asia, Europe-Americas differ significantly.


```
In [302]: options(repr.plot.width=4, repr.plot.height=4)

plot(mod, 1)
```




```
In [303]: leveneTest(income ~ region - 1, data = infmort, center=mean)
```

	Df	F value	Pr(>F)
group	3	5.869245	0.0009796424
	101	NA	NA

- In the residuals vs fitted plot, the errors seem to have different than zero mean and non-constant variances. Also, from leveneTest, the p value is less than 5% α level, hence we reject the null hypothesis.

```
In [305]: plot(mod, 2)
```



```
In [304]: # Shapiro-Wilk test  
shapiro.test(x = mod$residuals)
```

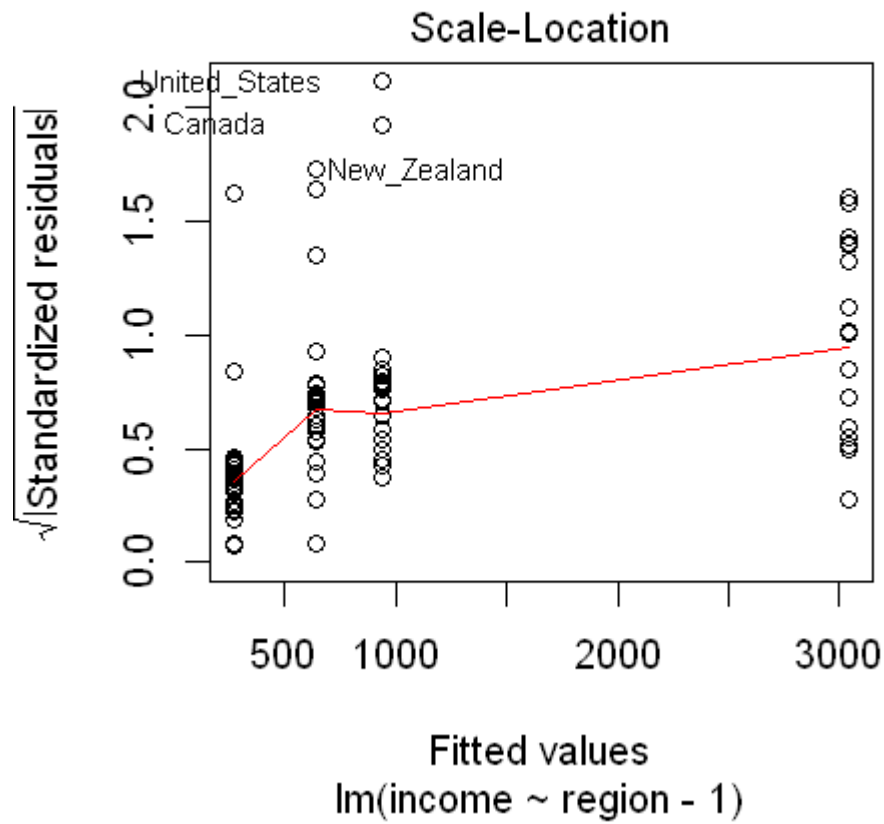
Shapiro-Wilk normality test

data: mod\$residuals
W = 0.75507, p-value = 6.061e-12

- Since p value for Shapro-Wilk test is less than 5% α level, normality assumption is not met. It


is clear from the normal Q-Q plot.

```
In [306]: plot(mod, 3)
```



- From the above plot, we can see there are few data points that are far away from residuals mean.

```
In [307]: plot(mod, 4)
```



```
In [309]: max(cooks.distance(mod))
```

0.225321798365331

- Maximum Cook's distance is ~0.22. Hence, we conclude there are no outliers.

(b) In case income varies with region, determine which pairs of regions are different.

```
In [310]: confint(mod)
```

	2.5 %	97.5 %
regionAfrica	-84.79121	631.2618
regionEurope	2548.16186	3532.2826
regionAsia	257.71835	1020.0150
regionAmericas	504.56751	1375.1716

```
In [311]: # Pairwise t-test with Bonferroni Correction
pairwise.t.test(infmort$income, infmort$region, p.adjust.method = "bonferroni")
```

Pairwise comparisons using t tests with pooled SD

data: infmort\$income and infmort\$region

	Africa	Europe	Asia
Europe	7.7e-14	-	-
Asia	1.00	7.2e-11	-
Americas	0.13	3.9e-08	1.00

P value adjustment method: bonferroni

```
In [312]: # Compute the analysis of variance
res.aov <- aov(income ~ region-1, data = infmort)
# Summary of the analysis
summary(res.aov)
TukeyHSD(res.aov)
```


```
              Df    Sum Sq Mean Sq F value Pr(>F)
region          4 201473156 50368289  45.48 <2e-16 ***
Residuals    101 111857493  1107500
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Tukey multiple comparisons of means
95% family-wise confidence level

Fit: aov(formula = income ~ region - 1, data = infmort)

\$region		diff	lwr	upr	p adj
Europe-Africa	2766.9869	1965.63247	3568.341	0.0000000	
Asia-Africa	365.6314	-323.00285	1054.266	0.5103789	
Americas-Africa	666.6343	-75.58589	1408.854	0.0945000	
Asia-Europe	-2401.3556	-3220.99422	-1581.717	0.0000000	
Americas-Europe	-2100.3527	-2965.50075	-1235.205	0.0000000	
Americas-Asia	301.0029	-460.92186	1062.928	0.7311095	

```
In [313]: options(repr.plot.width=6, repr.plot.height=3)
plot(TukeyHSD(res.aov), las=1, cex.axis = 0.5, col = 'blue')
```




Here, we can see that Africa-Europe, Europe-Asia, Europe-Americas differ significantly.

(c) In case you need a transformation for the response, re-fit the model and make a comparison with the previous results.

- We can try log transformation

```
In [385]: options(repr.plot.width=6, repr.plot.height=3)
ggplot(infmort, aes(x = region, y = log(income), fill=region)) +
  geom_boxplot()
```



```
In [386]: mod = lm(log(income) ~ region-1, data = infmort)
summary(mod)
anova(mod)
```

Call:

```
lm(formula = log(income) ~ region - 1, data = infmort)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-1.85632 -0.68738 -0.08462  0.41350  2.92786
```

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)
regionAfrica    5.0818     0.1630   31.18  <2e-16 ***
regionEurope    7.8627     0.2240   35.10  <2e-16 ***
regionAsia      5.6619     0.1735   32.63  <2e-16 ***
regionAmericas  6.3404     0.1982   32.00  <2e-16 ***
---

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9503 on 101 degrees of freedom

Multiple R-squared: 0.977, Adjusted R-squared: 0.9761

F-statistic: 1073 on 4 and 101 DF, p-value: < 2.2e-16

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
region	4	3877.16635	969.2915882	1073.336	8.994771e-82
Residuals	101	91.20952	0.9030646	NA	NA

```
In [387]: R_squared = var(mod$fitted.values)/var(log(infmort$income))
R_squared
```

```
0.515925905112566
```

```
In [388]: confint(mod)
```

	2.5 %	97.5 %
regionAfrica	4.758536	5.405132
regionEurope	7.418343	8.307003
regionAsia	5.317673	6.006027
regionAmericas	5.947370	6.733526

```
In [389]: # Pairwise t-test with Bonferroni Correction
pairwise.t.test(log(infmort$income),infmort$region,p.adjust.method = "bonferroni")
```

Pairwise comparisons using t tests with pooled SD

data: log(infmort\$income) and infmort\$region

	Africa	Europe	Asia
Europe	4.4e-16	-	-
Asia	0.099	4.1e-11	-
Americas	2.1e-05	1.0e-05	0.069

P value adjustment method: bonferroni

```
In [390]: # Compute the analysis of variance
res.aov <- aov(log(income) ~ region-1, data = infmort)
# Summary of the analysis
summary(res.aov)
TukeyHSD(res.aov)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
region	4	3877	969.3	1073	<2e-16 ***
Residuals	101	91	0.9		


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Tukey multiple comparisons of means
95% family-wise confidence level

Fit: aov(formula = log(income) ~ region - 1, data = infmort)

\$region		diff	lwr	upr	p adj
Europe-Africa		2.7808389	2.057216174	3.5044617	0.0000000
Asia-Africa		0.5800160	-0.041820426	1.2018524	0.0767083
Americas-Africa		1.2586139	0.588389434	1.9288384	0.0000210
Asia-Europe		-2.2008229	-2.940956307	-1.4606896	0.0000000
Americas-Europe		-1.5222250	-2.303453357	-0.7409966	0.0000098
Americas-Asia		0.6785979	-0.009419797	1.3666157	0.0546802

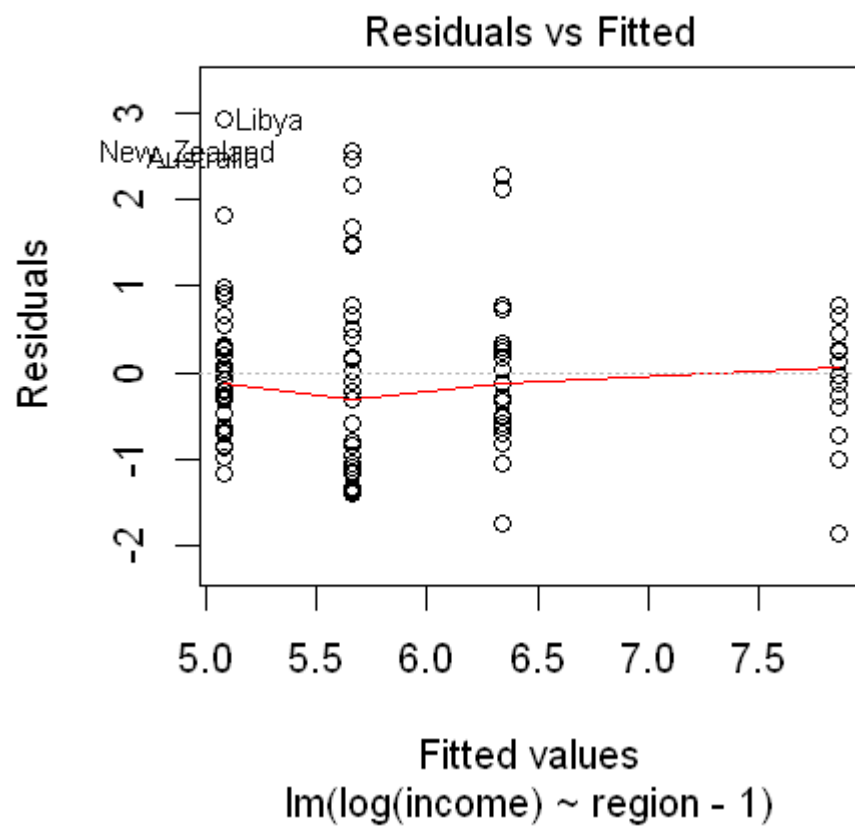

```
In [391]: options(repr.plot.width=6, repr.plot.height=3)
plot(TukeyHSD(res.aov), las=1, cex.axis = 0.5, col = 'blue')
```



Here, we can see that Africa-Europe, Europe-Asia, Europe-Americas, and Americas-Africa differ significantly. Asia-Africa and Americas-Asia have p-value close to 5% α level.

```
In [392]: options(repr.plot.width=4, repr.plot.height=4)
```

```
plot(mod, 1)
```




```
In [393]: leveneTest(log(income) ~ region - 1, data = infmort, center=mean)
```

	Df	F value	Pr(>F)
group	3	3.32169	0.02280456
	101	NA	NA

- In the residuals vs fitted plot, the errors seem to have different than zero mean and non-constant variances. Also, from leveneTest, the p value is less than 5% α level, hence we reject the null hypothesis.

```
In [394]: plot(mod, 2)
```



```
In [395]: # Shapiro-Wilk test
shapiro.test(x = mod$residuals)
```

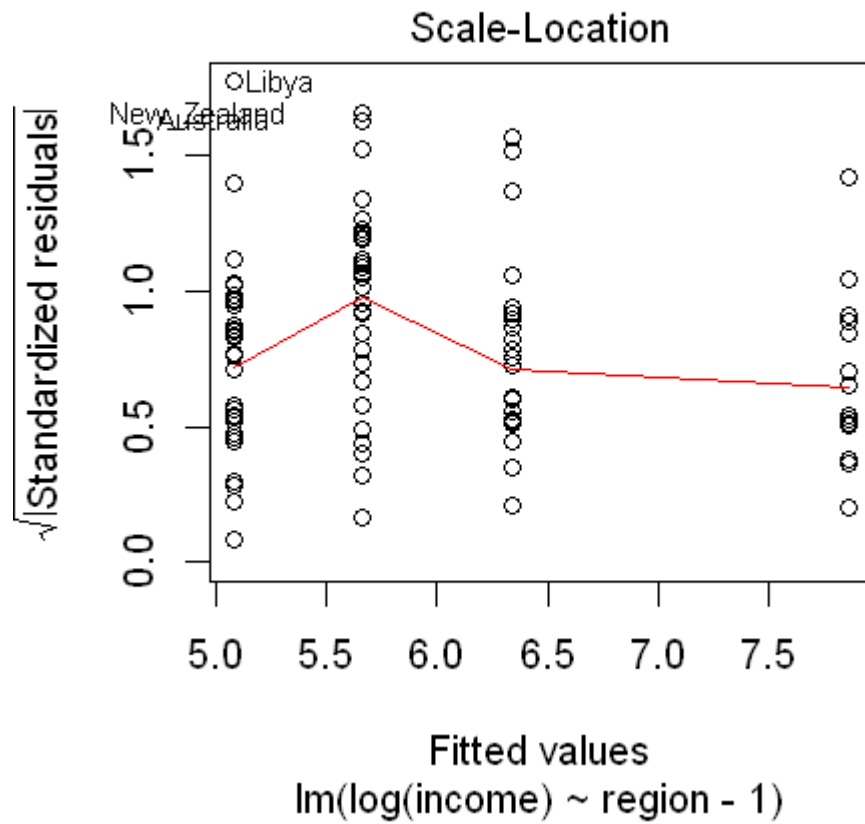
Shapiro-Wilk normality test

data: mod\$residuals
W = 0.94863, p-value = 0.0004753

- Since p value for Shapro-Wilk test is less than 5% α level, normality assumption is not met. It


is clear from the normal Q-Q plot.

```
In [396]: plot(mod, 3)
```



- From the above plot, we can see there are few data points that are far away from residuals mean.

```
In [397]: plot(mod, 4)
```



```
In [398]: max(cooks.distance(mod))
```

0.074092327884585

- Maximum Cook's distance is ~0.07. Hence, we conclude there are no outliers.

- After transforming income with log, we can see distinction between few more region was more clear.

In []: